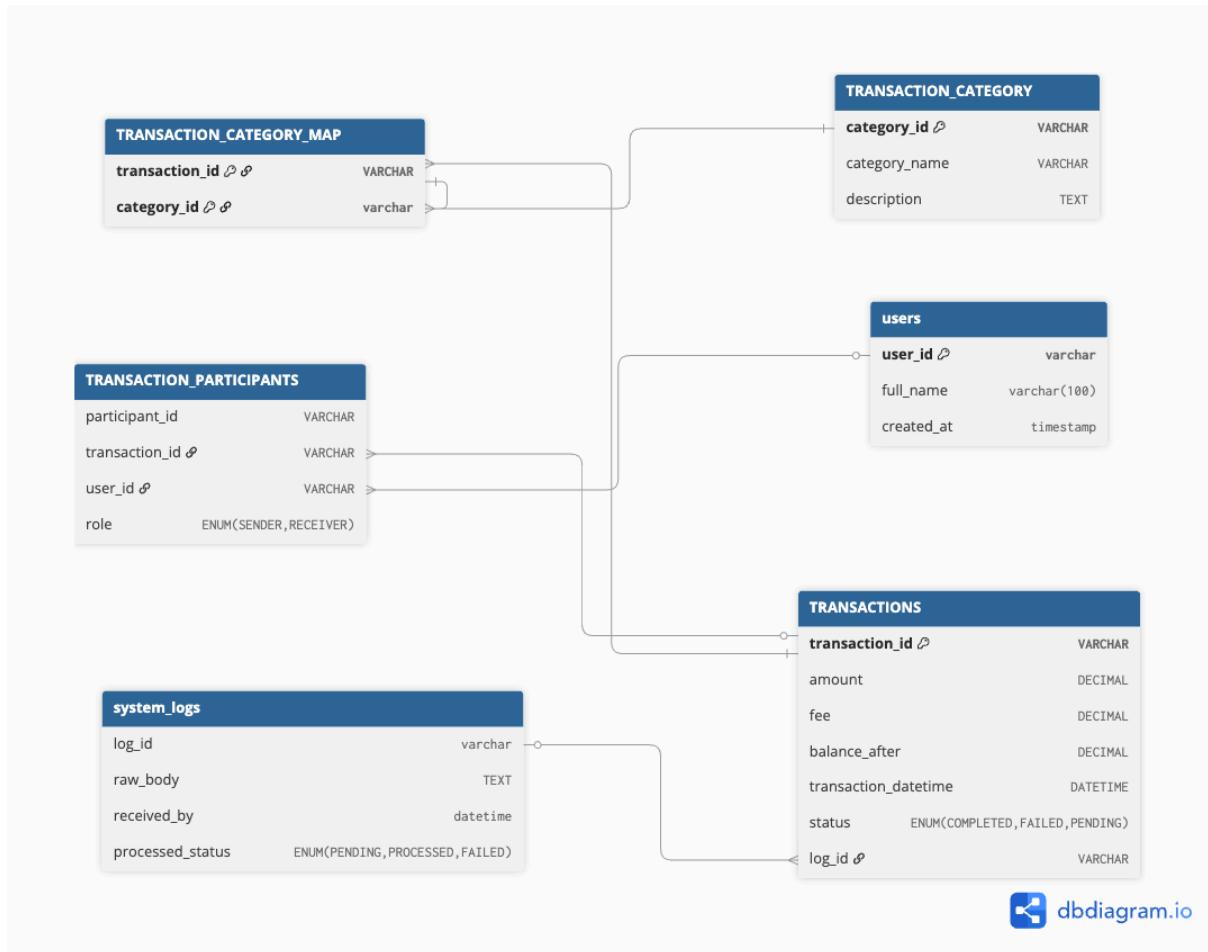


Database Design Document

Entity Relation diagram (ERD)



Design rationale and justification

- Why we separated the raw logs and transaction data

The design maintains a clear separation between raw SMS data (**SYSTEM_LOGS**) and processed transaction data (**TRANSACTIONS**). This architectural decision provides several benefits:

1. Complete audit trail - raw messages are never modified
2. Reprocessing capability - failed transactions can be reprocessed from original SMS

- 3. Data integrity - structured transaction data is validated separately
- 4. System reliability - processing failures don't lose original data
- The 1:1 relationship between SYSTEM_LOGS and TRANSACTIONS ensures that each

SMS can only generate one transaction, preventing duplicate processing while maintaining the ability to trace any transaction back to its source message.
- MANY-TO-MANY CATEGORY RELATIONSHIP

The M:N relationship between TRANSACTIONS and TRANSACTION_CATEGORY through the TRANSACTION_CATEGORY_MAP junction table allows flexible categorization.

BENEFITS:

- A single transaction can belong to multiple categories (e.g., a bank deposit that also involves a user payment)
- Sophisticated reporting and analysis without data duplication
- Easy addition of new categories without schema changes
- Flexible business logic for transaction classification

The junction table uses a composite primary key (transaction_id, category_id) to prevent duplicate category assignments while maintaining referential integrity through foreign keys

- PARTICIPANT TRACKING DESIGN

The TRANSACTION_PARTICIPANTS table implements a role-based model where each transaction records who was involved and their role (SENDER/RECEIVER).

This design supports:

1. Complete transaction history per user
2. Bilateral transaction tracking (sender and receiver perspectives)
3. Future extensibility for multi-party transactions

4. Clear accountability and reporting capabilities
5. Separation of concerns (users vs. their participation in transactions)

The participant_id primary key allows the same user to have multiple roles in different transactions while maintaining unique participant records.

DATA TYPE AND CONSTRAINT CHOICES

DECIMAL(12,2) for Financial Amounts:

- Prevents floating-point errors in monetary calculations
- Supports amounts up to 999,999,999,999.99 RWF
- Exact precision for accounting requirements

VARCHAR(20) for IDs:

- Human-readable identifiers from SMS messages
- Maintains uniqueness while being memorable
- Sufficient length for alphanumeric transaction IDs

ENUM Types for Status Fields:

- Ensures data consistency (only valid values accepted)
- Reduces storage compared to VARCHAR
- Self-documenting (valid states are explicit)
- Database-level validation

CHECK Constraints:

- Enforces business rules at database level
- Prevents invalid data entry (negative amounts, etc.)
- Cannot be bypassed by application code
- Provides automatic validation

Data dictionary

USERS Table

Column	Data_type	Constraints	Description
user_id	varchar(20)	Primary Key	Unique identifier
full_name	varchar(100)	NOT NULL	User full name
created_at	TIMESTAMP	Default current	Account creation

		timestamp	timestamp
--	--	-----------	-----------

SYSTEM_LOGS Table

Column	Data_type	Constraints	Description
log_id	varchar(20)	Primary Key	Unique Identifier
raw_body	TEXT	NOT NULL	Complete un processed sms
received_at	DATETIME	NOT NULL	Time stamp when the sms was received
processed_status	ENUM	DEFAULT PENDING	Processed status : PENDING , PROCESSED, FAILED

TRANSACTIONS Table

Column	Data_type	Constraints	Description
transaction_id	varchar(20)	Primary Key	Unique identifier
amount	decimal(12,2)	NOT NULL < check > 0	Transaction amount in RWF
fee	decimal(10,2)	DEFAULT 0 CHECK >=0	Fee charged for transaction
balance_after	decimal(12,2)		Account balance after transaction
transaction_datetime	Datetime	NOT NULL	When transaction occured
status	ENUM		Transaction status : COMPLETED , FAILED ,

			PENDING
log_id	varchar(20)	Foreign Key	Foreign key to sms log

TRANSACTION_CATEGORY Table

Column	Data_type	Constraints	Description
category_id	varchar(20)	Primary Key	Unique identifier
category_name	varchar(50)	Unique not null	Display name of the category
description	TEXT	null	description

TRANSACTION_CATEGORY_MAP Table

Column	Data_type	Constraints	Description
transaction_id	varchar(20)	Primary key , Foreign key	Foraging key to the transaction table
category_id	varchar(20)	Primary key , Foreign key	Foreign key to the transaction_category table

TRANSACTION PARTICIPANTS

Column	Data_type	Constraints	Description
participant_id	varchar(20)	Primary key	Unique identifier
transaction_id	varchar(20)	Foreign key	Foreign key to the transaction table
user_id	varchar(20)	Foreign key	Foreign key to

			the users table
role	ENUM	NOT NULL	Participant role : receiver , sender

Sample queries

TRANSACTION HISTORY FOR A SPECIFIC USER

Query

```

SELECT
    t.transaction_id,
    t.amount,
    t.fee,
    t.balance_after,
    t.transaction_datetime,
    t.status,
    tp.role,
    u.full_name
FROM TRANSACTIONS t
JOIN TRANSACTION_PARTICIPANTS tp
    ON t.transaction_id = tp.transaction_id
JOIN USERS u
    ON tp.user_id = u.user_id
WHERE u.user_id = 'U102'
ORDER BY t.transaction_datetime DESC;

```

Result

```

mysql> SELECT
    ->     t.transaction_id,
    ->     t.amount,
    ->     t.fee,
    ->     t.balance_after,
    ->     t.transaction_datetime,
    ->     t.status,
    ->     tp.role,
    ->     u.full_name
    -> FROM TRANSACTIONS t
    -> JOIN TRANSACTION_PARTICIPANTS tp
    ->     ON t.transaction_id = tp.transaction_id
    -> JOIN USERS u
    ->     ON tp.user_id = u.user_id
    -> WHERE u.user_id = 'U102'
    -> ORDER BY t.transaction_datetime DESC;
+-----+-----+-----+-----+-----+-----+-----+-----+
| transaction_id | amount | fee | balance_after | transaction_datetime | status | role | full_name |
+-----+-----+-----+-----+-----+-----+-----+-----+
| 17818959211 | 2000.00 | 0.00 | 38400.00 | 2024-05-11 18:48:42 | COMPLETED | RECEIVER | Samuel Carter |
| 51732411227 | 600.00 | 0.00 | 400.00 | 2024-05-10 21:32:32 | COMPLETED | RECEIVER | Samuel Carter |
+-----+-----+-----+-----+-----+-----+-----+-----+
2 rows in set (0.089 sec)

```

Explanation

This query joins three tables to provide complete transaction context for user U102 (Samuel Carter), showing their role in each transaction.

TRANSACTIONS BY CATEGORY

Query

SELECT

```

tc.category_name,
COUNT(tcm.transaction_id) as transaction_count,
SUM(t.amount) as total_amount,
AVG(t.amount) as average_amount
FROM TRANSACTION_CATEGORY tc
JOIN TRANSACTION_CATEGORY_MAP tcm
    ON tc.category_id = tcm.category_id
JOIN TRANSACTIONS t
    ON tcm.transaction_id = t.transaction_id
WHERE t.status = 'COMPLETED'
GROUP BY tc.category_name
ORDER BY total_amount DESC;

```

Result

```

mysql> SELECT
->     tc.category_name,
->     COUNT(tcm.transaction_id) as transaction_count,
->     SUM(t.amount) as total_amount,
->     AVG(t.amount) as average_amount
->   FROM TRANSACTION_CATEGORY tc
->   JOIN TRANSACTION_CATEGORY_MAP tcm
->     ON tc.category_id = tcm.category_id
->   JOIN TRANSACTIONS t
->     ON tcm.transaction_id = t.transaction_id
-> WHERE t.status = 'COMPLETED'
-> GROUP BY tc.category_name
-> ORDER BY total_amount DESC;
+-----+-----+-----+-----+
| category_name | transaction_count | total_amount | average_amount |
+-----+-----+-----+-----+
| Bank Deposit      |           1 |    40000.00 | 40000.000000 |
| Payment Sent       |           3 |     3600.00 | 1200.000000 |
| Payment Received    |           1 |     2000.00 | 2000.000000 |
+-----+-----+-----+-----+
3 rows in set (0.063 sec)

```

Explanation

Provides financial summary by transaction category, useful for business intelligence and reporting.

UNPROCESSED SMS LOGS

Query

```

SELECT
  sl.log_id,
  sl.raw_body,
  sl.received_at,
  sl.processed_status
FROM SYSTEM_LOGS sl
LEFT JOIN TRANSACTIONS t
  ON sl.log_id = t.log_id
WHERE sl.processed_status != 'PROCESSED'
  OR t.transaction_id IS NULL
ORDER BY sl.received_at ASC;

```

Result

```
mysql> SELECT
    ->     sl.log_id,
    ->     sl.raw_body,
    ->     sl.received_at,
    ->     sl.processed_status
    -> FROM SYSTEM_LOGS sl
    -> LEFT JOIN TRANSACTIONS t
    ->     ON sl.log_id = t.log_id
    -> WHERE sl.processed_status != 'PROCESSED'
    ->     OR t.transaction_id IS NULL
    -> ORDER BY sl.received_at ASC;
Empty set (0.011 sec)
```

Explanation

Critical for monitoring system health and identifying processing failures or backlogs. And for this database (No unprocessed logs in current dataset - all SMS messages have been successfully processed)

DAILY TRANSACTION SUMMARY

Query

```
SELECT
    DATE(transaction_datetime) as transaction_date,
    COUNT(*) as total_transactions,
    SUM(amount) as total_volume,
    AVG(amount) as average_transaction,
    SUM(fee) as total_fees_collected
FROM TRANSACTIONS
WHERE status = 'COMPLETED'
GROUP BY DATE(transaction_datetime)
ORDER BY transaction_date DESC;
```

Result

```

mysql> SELECT
->     DATE(transaction_datetime) as transaction_date,
->     COUNT(*) as total_transactions,
->     SUM(amount) as total_volume,
->     AVG(amount) as average_transaction,
->     SUM(fee) as total_fees_collected
-> FROM TRANSACTIONS
-> WHERE status = 'COMPLETED'
-> GROUP BY DATE(transaction_datetime)
-> ORDER BY transaction_date DESC;
+-----+-----+-----+-----+-----+
| transaction_date | total_transactions | total_volume | average_transaction | total_fees_collected |
+-----+-----+-----+-----+-----+
| 2024-05-11      | 2                | 42000.00    | 21000.000000        | 0.00                 |
| 2024-05-10      | 3                | 3600.00     | 1200.000000         | 0.00                 |
+-----+-----+-----+-----+-----+
2 rows in set (0.008 sec)

```

Explanation

Enables daily business reporting and trend identification for decision-making.

Security and data integrity rules

CHECK CONSTRAINTS

Check constraints enforce business rules at the database level, preventing invalid data entry regardless of application logic.

Query

```

INSERT INTO TRANSACTIONS (transaction_id, amount, fee, balance_after,
transaction_datetime, status, log_id) VALUES
('D002', -40000, 0, 40400, '2024-05-11 18:43:49', 'COMPLETED', 'L1004');

```

Result

```

mysql> INSERT INTO TRANSACTIONS (transaction_id, amount, fee, balance_after, transaction_datetime, status, log_id) VALUES
-> ('D002', -40000, 0, 40400, '2024-05-11 18:43:49', 'COMPLETED', 'L1004');
ERROR 3819 (HY000): Check constraint 'transactions_chk_1' is violated.
mysql> ■

```

Explanation

As you can see above when we try to insert a transaction record with negative(-) amount the check constraints are violated the same thing happens when the amount is equal to 0 . this ensures that no invalid amount is ever in the table

UNIQUE CONSTRAINTS

Unique constraints prevent duplicate data and enforce one-to-one relationships.

Example : all the primary keys are unique and don't allow duplicate values

INDEXES

Indexes allow for fast searching of records in tables

Query

```
EXPLAIN SELECT * FROM TRANSACTIONS WHERE status =  
'COMPLETED'  
ORDER BY transaction_datetime DESC;
```

Result

```
mysql> EXPLAIN SELECT * FROM TRANSACTIONS WHERE status = 'COMPLETED'  
-> ORDER BY transaction_datetime DESC;  
+-----+  
| EXPLAIN  
|  
+-----+  
| -> Sort: transactions.transaction_datetime DESC (cost=1 rows=5)  
|   -> Index lookup on TRANSACTIONS using idx_tx_status (status = 'COMPLETED'), with index condition: (transactions.`status` = 'COMPLETED') (cost=1 rows=5)  
|  
+-----+  
1 row in set (0.008 sec)
```

ADDITIONAL SECURITY MEASURES

- NOT NULL CONSTRAINTS: Critical fields (amounts, names, dates) cannot be null, preventing incomplete records that could compromise data integrity.
- DEFAULT VALUES: Automatic timestamp generation for created_at ensures proper audit trails without relying on application logic.

- DECIMAL PRECISION: Financial amounts use DECIMAL(12,2) to prevent rounding errors inherent in FLOAT types. This is critical for accounting accuracy and financial compliance.
- IMMUTABLE LOGS: SYSTEM_LOGS table design encourages append-only operations, preserving complete audit history. Raw SMS data is never modified or deleted, maintaining compliance with financial regulations.
- COMPOSITE KEYS: Junction table (TRANSACTION_CATEGORY_MAP) uses composite primary key (transaction_id, category_id) preventing duplicate mappings while maintaining referential integrity.
- VARCHAR LENGTH LIMITS: All VARCHAR columns have appropriate length limits to prevent buffer overflow attacks and ensure data quality.
- TIMESTAMP TRACKING: All major tables include timestamp fields (created_at, received_at, transaction_datetime) providing complete temporal audit trail.